



# Overview of GPU usage on Alpine

# GPU nodes on Alpine

---




12 NVIDIA GPU NODES (3  
GPU EACH)



8 AMD GPU NODES (3 GPU  
EACH)

# GPU partition on Alpine



AA100 -> NVIDIA RELATED  
PARTITION; MAX WALLTIME  
24 HOURS



AMI100 -> AMD RELATED  
PARTITION; MAX WALLTIME  
24 HOURS



ATESTING\_A100 -> NVIDIA  
TESTING PARTITION; MAX  
WALLTIME 1 HOUR



ATESTING\_AMI100 -> AMD  
TESTING PARTITION; MAX  
WALLTIME 1 HOUR



IMPORTANT: TO REQUEST  
MORE THAN 24 HOURS  
WALLTIME USE --QOS=LONG

# Access to NVIDIA Job partition

```
#!/bin/bash
```

```
#SBATCH --job-name=full_genome_nn_hp_opt_gpu  
#SBATCH --output=full_genome_nn_hp_opt_gpu.out  
#SBATCH --error=full_genome_nn_hp_opt_gpu.err  
#SBATCH --partition=aa100  
#SBATCH --nodes=1  
#SBATCH --ntasks=45  
#SBATCH --gres=gpu:1  
#SBATCH --account=amc-general  
#SBATCH --time=24:00:00
```

- To access an AMD partition add: **--partition=ami100**
- Here, the number of gpus are allocated with: **--gres=gpu:1**



# Watch out for cache size

- /home or /temp have a very small size. Please include the following in your slurm script.

```
export TMP=/scratch/alpine/$USER
```

```
export TEMP=/scratch/alpine/$USER
```

```
export TMPDIR=/scratch/alpine/$USER
```

```
export TEMPDIR=/scratch/alpine/$USER
```

# Access to gpu nodes testing partitions

- **[NVIDIA]** sinteractive --  
partition=atesting\_a100 --qos=testing --  
time=00:05:00 --gres=gpu:1 --ntasks=2
- **[AMD]** sinteractive --  
partition=atesting\_mi100 --qos=testing --  
time=00:05:00 --gres=gpu:1 --ntasks=2


# Slurm Quality of service (qos)

- Used to modify or constrain characteristics that a job can have.
- **--qos=normal** corresponds to a walltime of 24 hours and is the default.
- **--qos=long** corresponds to a walltime of up to 7 days
- **--qos=mem** corresponds to high memory jobs only (up to 1TB)



# Fairshare overview

- Difference between the portion of computing resource that has been promised and the amount of resources that has been consumed.
- Level fairshare of 1 indicates average priority compared to other users in that account (amc-general)
- **module load slurmttools; levelsf \$USER**





# Job priority calculation formula

```
Job_priority =  
    site_factor +  
    (PriorityWeightAge) * (age_factor) +  
    (PriorityWeightAssoc) * (assoc_factor) +  
    (PriorityWeightFairshare) * (fair-share_factor) +  
    (PriorityWeightJobSize) * (job_size_factor) +  
    (PriorityWeightPartition) * (partition_factor) +  
    (PriorityWeightQOS) * (QOS_factor) +  
    SUM(TRES_weight_cpu * TRES_factor_cpu,  
        TRES_weight_<type> * TRES_factor_<type>,  
        ...)  
    - nice_factor
```

# Package availability (1)

Some packages that have been built and accessible through lmod.

Adding new packages through lmod takes a lot of round of approval so it is recommended to build them locally.

Solutions: (cmake+make), Anaconda, pip, containers, spack etc ...

Submit a ticket at [rc-help@Colorado.edu](mailto:rc-help@Colorado.edu) so that I can build it for you locally.

# Package availability for ML (2)

- Cuda 11.2, Cuda 11.3 and Cuda 11.4 on Alpine.
- Only cudnn 8.1 and 8.2 on Alpine.
- Can be problematic for DL build with GPU compatibility

## GPU

Version	Python version	Compiler	Build tools	cuDNN	CUDA
tensorflow-2.13.0	3.8-3.11	Clang 16.0.0	Bazel 5.3.0	8.6	11.8
tensorflow-2.12.0	3.8-3.11	GCC 9.3.1	Bazel 5.3.0	8.6	11.8
tensorflow-2.11.0	3.7-3.10	GCC 9.3.1	Bazel 5.3.0	8.1	11.2
tensorflow-2.10.0	3.7-3.10	GCC 9.3.1	Bazel 5.1.1	8.1	11.2
tensorflow-2.9.0	3.7-3.10	GCC 9.3.1	Bazel 5.0.0	8.1	11.2
tensorflow-2.8.0	3.7-3.10	GCC 7.3.1	Bazel 4.2.1	8.1	11.2
tensorflow-2.7.0	3.7-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.6.0	3.6-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.5.0	3.6-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.4.0	3.6-3.8	GCC 7.3.1	Bazel 3.1.0	8.0	11.0
tensorflow-2.3.0	3.5-3.8	GCC 7.3.1	Bazel 3.1.0	7.6	10.1

# Pytorch installation on Alpine [NVIDIA] (1)

1. After you log into the Alpine cluster, please load the slurm modules and request allocation so that you can install the packages:

```
module load slurm/alpine  
acompile --ntasks=4
```



2. Load anaconda, create your environment with python 3.10 and activate it.

```
module load anaconda  
conda create -n pytorch_env python=3.10  
conda activate pytorch_env
```



3. Install pytorch, pytorch-cuda. You could also install torchvision and torchaudio if needed for your workflow.

```
conda install pytorch==2.0.0 torchvision==0.15.0 torchaudio==2.0.0 pytorch-cuda=11.8 -c pytorch -c nvidia
```



4. Install cuda-toolkit 11.8.0

```
conda install -c "nvidia/label/cuda-11.8.0" cuda-toolkit
```



# Pytorch installation on Alpine [NVIDIA] (2)

## 5. Install nvidia-cudnn 8.6.0

```
pip install nvidia-cudnn-cu11==8.6.0.163
```



## 6. To test that your installation is working you will need to exit "acompile" first and load on the NVIDIA gpu debug partition on Alpine"

```
conda deactivate  
exit  
sinteractive --partition=atesting_a100 --qos=testing --time=00:05:00 --gres=gpu:1 --ntasks=2  
module load anaconda  
conda activate pytorch_env  
python
```



## 7. Finally, run the following:

```
>>>import torch  
>>>print(torch.cuda.is_available())  
True
```



## 8. Make sure to exit the GPU debug node partition after testing the installation.

```
$ exit  
exit
```





# Tensorflow installation on Alpine [NVIDIA] (1)

1. After you log into the Alpine cluster, please load the Slurm modules and request allocation so that you can install the packages:

```
module load slurm/alpine  
acompile --ntasks=4 --time=01:30:00
```



2. Load anaconda, create your environment with python 3.9 and activate it.

```
module load anaconda  
conda create -n tf_env python=3.9  
conda activate tf_env
```



3. Install cudnn 8.6.0

```
pip install nvidia-cudnn-cu11==8.6.0.163
```



4. Install cuda-toolkit 11.8.0

```
conda install -c "nvidia/label/cuda-11.8.0" cuda-toolkit
```



# Tensorflow installation on Alpine [NVIDIA] (2)

5. Install Tensorflow 2.12.0. Also note that tensorflow, cuda and cudnn have to follow a strict versioning: <https://www.tensorflow.org/install/source#gpu>

```
python3 -m pip install tensorflow==2.12.0
```

6. Export the correct paths by following this guide here: <https://www.tensorflow.org/install/pip>

```
mkdir -p $CONDA_PREFIX/etc/conda/activate.d
echo 'CUDNN_PATH=$(dirname $(python -c "import nvidia.cudnn;print(nvidia.cudnn.__file__)"))' >> $CONDA_PREFIX/etc/conda/activate.d/activate.d
echo 'export LD_LIBRARY_PATH=$CONDA_PREFIX/lib:$CUDNN_PATH/lib:$LD_LIBRARY_PATH' >> $CONDA_PREFIX/etc/conda/activate.d/activate.d
source $CONDA_PREFIX/etc/conda/activate.d/env_vars.sh
export LD_LIBRARY_PATH=$CONDA_PREFIX/lib/python3.9/site-packages/nvidia/cudnn/lib:$LD_LIBRARY_PATH
export LD_LIBRARY_PATH=$CONDA_PREFIX/lib:$LD_LIBRARY_PATH
export PATH=$CONDA_PREFIX/bin:$PATH
export XLA_FLAGS=--xla_gpu_cuda_data_dir=$CONDA_PREFIX
```

7. Install Tensorrt and export PATH:

```
pip install nvidia-tensorrt==8.4.1.5
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/python3.9/site-packages/tensorrt
```

# Tensorflow installation on Alpine [NVIDIA] (3)

8. We link libnvinfer.so.8 to libnvinfer.so.7 :

```
ln -s $CONDA_PREFIX/lib/python3.9/site-packages/tensorrt/libnvinfer.so.8 $CONDA_PREFIX/lib/python3.9/site-packages/t  
ln -s $CONDA_PREFIX/lib/python3.9/site-packages/tensorrt/libnvinfer_plugin.so.8 $CONDA_PREFIX/lib/python3.9/site-pa
```

9. To test that your installation is working you will need to exit "acompile" first and load on the NVIDIA GPU debug partition on Alpine"

```
conda deactivate  
exit  
sinteractive --partition=atesting_a100 --qos=testing --time=00:05:00 --gres=gpu:1 --ntasks=2  
module load anaconda  
conda activate tf_env  
python3 -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
```

10. Make sure to exit the GPU debug node partition after testing the installation.

```
$ exit  
exit
```

# Containers (1)

- Apptainer/Singularity can now be built directly on the cluster
- Can be built either from a definition file or converted from a docker image.
- e.g: `apptainer build -f R_spack_env.sif R_spack_env.def`



# Containers (2)

- `module load singularity`
- `export ALPINE_SCRATCH=/gpfs/alpine1/scratch/$USER`
- `export SINGULARITY_TMPDIR=$ALPINE_SCRATCH/singularity/tmp`
- `export SINGULARITY_CACHEDIR=$ALPINE_SCRATCH/singularity/cache`  
`mkdir -pv $SINGULARITY_CACHEDIR $SINGULARITY_TMPDIR`